

Testing 1, 2, 3...

Christopher R. Genovese

Department of Statistics & Data Science

19 Sep 2024
Session #8

Plan

Motivation

Plan

Motivation

Practicalities

Plan

Motivation

Practicalities

Practice

Plan

Motivation

Practicalities

Practice

Activity

Announcements

- Konrad will be announcing office hours; keep an eye out for them.
- **Homework:** `state-machines` assignment is up, due Thursday 26 Sep.

Plan

Motivation

Practicalities

Practice

Activity

Complexity can have consequences!

Bugs will happen:

- the crash of the [Mars Climate Orbiter](#) (1998),
- a [failure of the national telephone network](#) (1990),
- a deadly medical device ([1985](#), 2000),
- a [massive Northeastern blackout](#) (2003),
- the [Heartbleed](#), [Goto Fail](#), [Shellshock](#) exploits (2012–2014),
- a 15-year-old [fMRI analysis software bug](#) that inflated significance levels (2015).

It is hard to know whether a piece of software is actually doing what it is supposed to do. It is easy to write a thousand lines of research code, then discover that your results have been wrong for months.

Discipline, design, and careful thought are all helpful in producing working software. But even more important is [effective testing](#), and that is the central topic for today.

Quick Testing Terminology

- **test** – a function that runs other code in a library or application codebase and checks its results for correctness
- **test suite** – a collection of tests on a related theme
- **unit** – (loosely) a piece of code with a small, well-defined purpose and scope.
- **assertion** – a claim about the state of a program or the result of a test
- **unit test** – an automated test of a unit, usually checking the result with several different inputs or configurations
- **property** – a logical invariant that a piece of code should satisfy
- **property test** – a test of a property that generates consistent inputs and upon failure, attempts to produce a near-minimal example that causes failure
- **fixture** – a context or setting that needs to be setup before tests are run (and usually torn down afterwards); these often create fakes/stubs/mocks that simulate entities our code should interact with in reality

Unit Testing

A “unit” is a vaguely defined concept that is intended to represent a small, well-defined piece of code. A unit is usually a function, method, class, module, or small group of related classes.

A test is simply some code that calls the unit with some inputs and checks that its answer matches an expected output.

Unit testing consists of writing tests that are

- focused on a small, low-level piece of code (a [unit](#))
- typically written by the programmer with standard tools
- fast to run (so can be run often, i.e. before every commit).

The benefits of unit testing are many, including

- Exposing problems early
- Making it easy to change (refactor) code without forgetting pieces or breaking things
- Simplifying integration of components
- Providing natural documentation of what the code should do
- Driving the design of new code.

Property Testing (aka Generative Testing)

A powerful technique for automatically testing logical invariants without having to create many small examples by hand. A **property** is a claim made about a specified set of quantities; we specify the set through **generators** that produce values of particular types and shapes.

Generators may be combined to form new generators that are more specific to our needs.

Property Testing (aka Generative Testing)

```
(def sort-idempotent-prop
  (for-all [v (gen/vector gen/int)]
    (= (sort v) (sort (sort v)))))

(quick-check 100 sort-idempotent-prop)
;; => {:result true,
;;     :pass? true,
;;     :num-tests 100,
;;     :time-elapsed-ms 28,
;;     :seed 1528580707376}
```

Property Testing (aka Generative Testing)

```
from hypothesis import given
from hypothesis.strategies import text

@given(text())
def test_decode_inverts_encode(s):
    assert decode(encode(s)) == s
```

Property Testing (aka Generative Testing)

One of the most useful features of Property/Generative testing is that on failures, the generators will /search/ for a (roughly) minimal/simple example that fails, to make it easier to identify the problem.

```
(def sorted-first-less-than-last-prop
  (for-all [v (gen/not-empty (gen/vector gen/int))]
    (let [s (sort v)]
      (< (first s) (last s)))))

(quick-check 100 sorted-first-less-than-last-prop)
;; => {:num-tests 5, :seed 1528580863556,
;;     :fail [[-3]],
;;     :failed-after-ms 1,
;;     :result false,
;;     :result-data nil,
;;     :failing-size 4,
;;     :pass? false,
;;     :shrunk
;;     {:total-nodes-visited 5, :depth 2,
;;      :pass? false, :result false, :result-data nil,
;;      :time-shrinking-ms 1,
;;      :smallest [[0]]}}
```

Property Testing: Generators/Strategies

```
from hypothesis.strategies import characters, text
from hypothesis import given
from unicodedata import category

names = text(
    characters(max_codepoint=1000, blacklist_categories=('Cc', 'Cs')),
    min_size=1).map(lambda s: s.strip()).filter(lambda s: len(s) > 0)

@given(names)
def test_names_match_our_requirements(name):
    assert len(name) > 0
    assert name == name.strip()
    for c in name:
        assert 1 <= ord(c) <= 1000
        assert category(c) not in ('Cc', 'Cs')
```

Property Testing: Generators/Strategies

```
from hypothesis import given
from hypothesis.extra.datetime import datetimes

project_date = datetimes(timezones=('UTC',), min_year=2000, max_year=2100)

@given(project_date)
def test_dates_are_in_the_right_range(date):
    assert 2000 <= date.year <= 2100
    assert date.tzinfo._tzname == 'UTC'
```


Property Testing: Generators/Strategies

```
test_that( "Reverse of reverse is identity",  
  forall( gen.c( gen.element(1:100) ), function(xs) expect_equal(rev(rev(xs)), xs))  
)  
  
test_that( "a is less than b + 1",  
  forall(list(a = gen.element(1:100), b = gen.unif(1,100, shrink.median = F))  
    , function(a, b) expect_lt( a, b + 1 ))
```

Property Testing: Generators/Strategies

```
test_that( "Reversed of concatenation is flipped concatenation of reversed",
  forall( list( as = gen.c( gen.element(1:100) )
               , bs = gen.c( gen.element(1:100) )
               , function(as,bs) expect_equal ( rev(c(as, bs)), c(rev(bs), rev(as)))
            )
  )
)
```

```
gen.df.of <- function ( n )
  gen.with (
    list( as = gen.c(of = n, gen.element(1:10) )
          , bs = gen.c(of = n, gen.element(10:20) )
        )
    , as.data.frame
  )
)
```

```
test_that( "Number of rows is 5",
  forall( gen.df.of(5), function(df) expect_equal(nrow(df), 5))
)
```

Plan

Motivation

Practicalities

Practice

Activity

Recommended Libraries: Python

`pytest` (unit testing+) and `Hypothesis` (generative testing)

```
from __future__ import annotations
import pytest
# ...
```

```
def test_kinds_factories():
    "Builtin kind factories"
    a = symbol('a')
```

```
assert constant(1).values == {1}
assert constant((2,)).values == {2}
assert constant((2, 3)).values == {vec_tuple(2, 3)}
```

```

assert either(0, 1).values == {0, 1}
assert weights_of(either(0, 1, 2).weights) == pytest.approx([as_quantity('2/3'),
                                                                as_quantity('1/3')])
assert lmap(str, values_of(either(a, 2 * a, 2).weights)) == ['<a>', '<2 a>']

```

...

```
with pytest.raises(KindError):
    k0 >> me1
```

Recommended Libraries: R

`testthat` (unit testing) and `hedgehog` (property testing) or `quickcheck` (property testing)

```
library(testthat)
```

```
source("foobar.R")
```

```
test_that("foo values are correct", {  
  expect_equal(foo(4), 8)  
  expect_equal(foo(2.2), 1.9)  
})
```

```
test_that("bar has correct limits", {  
  expect_lt(bar(4, c(1, 90)), option = TRUE, 8)  
})
```

```
test_that("bar throws an error on bad inputs", {  
  expect_error(bar(-4, c(1, 10))) # test passes if bar calls stop()  
                                 # or throws an error here  
})
```

Test Structure

In these frameworks, *tests* are functions flagged by a special name, e.g., `test_`.

Each test contains one or more *assertions*. These can take different forms. The `test_that` assertions are functions like `expect_equal` or `expect_error`.

They take the result to be tested and (usually) the correct result and compare them in a suitable way.

Each test can have many assertions; the failure of any one causes the test to fail.

In practice, you can select which tests are run, though we usually just run them all.

Tips and Practices

- Tests are commonly kept in separate source files from the rest of your code. In a long-running project, you may have a `test/` folder containing test code for each piece of your project, plus any data files or other bits needed for the tests.
- All tests can now be run with a single command (e.g. using `testthat`'s `test_dir` function or Python's `pytest` module)
- Run tests **often**. It is common to set up a *hook* that runs your tests before each commit, and perhaps rejects the commit if the tests fail.
- Every time you check your code, such as at the repl or with an example run, *make a test out of it*. Every time you encounter a bug or other failure, make a test out of it. Every example you put in your documentation can produce a test.
- There is a wide variety of built-in assertions in common testing libraries, including for instance asserting that a piece of code throws an error. There may also be third party libraries that add additional assertions and tools; these can be included as a “dev dependency” without affecting your users. And you can always write your own if you need something special repeatedly.
- It is valuable when possible to write some tests **before** you implement a function. This can help you understand (and even document) what the function needs to do, including edge cases.
- Make tests **replicable**: If a test involves random data, what do you do when the test fails? You need some way to know what random values it used so you can figure out why the test fails.

Plan

Motivation

Practicalities

Practice

Activity

Rapid-Fire Testing 1

Scenario. Find the maximum sum of a subsequence

Function name: `max_sub_sum(arr)`

Write a function that takes as input a vector of n numbers and returns the maximum sum found in any *contiguous* subvector of the input. (We take the sum of an empty subvector to be zero.)

For example, in the vector $[1, -4, 4, 2, -2, 5]$, the maximum sum is 9, for the subvector $[4, 2, -2, 5]$.

As we've seen, there's a clever algorithm for doing this fast, with an interesting history related to our department. But that's not important right now. How do we test it?

(If you want to implement it – there's a repository problem for that! Try the `max-sub-sum` exercise.)

Test ideas? Consider properties!

Rapid-Fire Testing 2

Scenario. Create a half-space function for a given vector.

Function name: `half_space_of(point)`

Given a vector in Euclidean space, return a boolean *function* that tests whether a *new* point is in the positive half space of the original vector. (The vector defines a perpendicular plane through the origin which splits the space in two: the positive half space and the negative half space.) An example:

```
foo <- half_space_of(c(2, 2))  
foo(c(1, 1)) == TRUE
```

Test ideas?

Plan

Motivation

Practicalities

Practice

Activity

Activity: Writing Tests for StateMachine

Write some tests for your state machine code. You can use the specification in the problem bank, if you don't remember.

When we start with tests, it's OK that those tests fail initially. so concentrate on the gist rather than the details.

THE END